

# PHASE 2 PROJECT

## GROUP MEMBERS

1. SHARON KALIKU
2. PAUL KAMAU
3. KIPKOSGEI KIPTUI
4. EZRA KIPCHIRCHIR
5. HERI KIMOTHO

## BUSINESS STAKEHOLDER

The real estate agency

## BUSINESS PROBLEM

A real estate agency wants to analyze the factors that influence the prices of houses in order to provide accurate pricing estimates to their clients. The agency aims to understand the relationship between various features of a house, such as the number of rooms, living area, basement area, overall quality, and other relevant factors, and how they affect the sale price.

The clients being, homeowners and potential house buyers have difficulty in making informed decisions regarding property investments, to make this decision, understanding the factors influencing housing prices in a specific area is necessary.

## OBJECTIVES

### REAL ESTATE AGENCY

- To identify the locations with the highest sales prices.
- To identify how seasonal trends affect sales.
- To predict prices of houses depending on the features.

Importing the necessary libraries that will be used to perform analysis on our data

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import statsmodels.api as sm
from scipy import stats
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
```

```

from sklearn.preprocessing import PolynomialFeatures
from sklearn.preprocessing import StandardScaler
data = pd.read_csv('kc_house_data.csv')
areas = pd.read_csv('deliverylocations.csv')

```

Defining functions to load and view the data

*# loading data*

```

def description_data(data):
    data = pd.read_csv(data)
    print("\n.....Info:.....")
    print(data.info())
    print("\n.....Describe:.....")
    print(data.describe())
    print("\n.....Head:.....")
    print(data.head())

```

```
description_data('kc_house_data.csv')
```

```

.....Info:.....
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21597 entries, 0 to 21596
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                     21597 non-null  int64
1   date                   21597 non-null  object
2   price                  21597 non-null  float64
3   bedrooms               21597 non-null  int64
4   bathrooms              21597 non-null  float64
5   sqft_living            21597 non-null  int64
6   sqft_lot               21597 non-null  int64
7   floors                 21597 non-null  float64
8   waterfront             19221 non-null  object
9   view                   21534 non-null  object
10  condition              21597 non-null  object
11  grade                  21597 non-null  object
12  sqft_above             21597 non-null  int64
13  sqft_basement          21597 non-null  object
14  yr_built               21597 non-null  int64
15  yr_renovated           17755 non-null  float64
16  zipcode                21597 non-null  int64
17  lat                    21597 non-null  float64
18  long                   21597 non-null  float64
19  sqft_living15          21597 non-null  int64
20  sqft_lot15             21597 non-null  int64
dtypes: float64(6), int64(9), object(6)
memory usage: 3.5+ MB

```

None

.....Describe:.....

	id	price	bedrooms	bathrooms
sqft_living \				
count	2.159700e+04	2.159700e+04	21597.000000	21597.000000
21597.000000				
mean	4.580474e+09	5.402966e+05	3.373200	2.115826
2080.321850				
std	2.876736e+09	3.673681e+05	0.926299	0.768984
918.106125				
min	1.000102e+06	7.800000e+04	1.000000	0.500000
370.000000				
25%	2.123049e+09	3.220000e+05	3.000000	1.750000
1430.000000				
50%	3.904930e+09	4.500000e+05	3.000000	2.250000
1910.000000				
75%	7.308900e+09	6.450000e+05	4.000000	2.500000
2550.000000				
max	9.900000e+09	7.700000e+06	33.000000	8.000000
13540.000000				

	sqft_lot	floors	sqft_above	yr_built
yr_renovated \				
count	2.159700e+04	21597.000000	21597.000000	21597.000000
17755.000000				
mean	1.509941e+04	1.494096	1788.596842	1970.999676
83.636778				
std	4.141264e+04	0.539683	827.759761	29.375234
399.946414				
min	5.200000e+02	1.000000	370.000000	1900.000000
0.000000				
25%	5.040000e+03	1.000000	1190.000000	1951.000000
0.000000				
50%	7.618000e+03	1.500000	1560.000000	1975.000000
0.000000				
75%	1.068500e+04	2.000000	2210.000000	1997.000000
0.000000				
max	1.651359e+06	3.500000	9410.000000	2015.000000
2015.000000				

	zipcode	lat	long	sqft_living15
sqft_lot15				
count	21597.000000	21597.000000	21597.000000	21597.000000
21597.000000				
mean	98077.951845	47.560093	-122.213982	1986.620318
12758.283512				
std	53.513072	0.138552	0.140724	685.230472
27274.441950				
min	98001.000000	47.155900	-122.519000	399.000000

```

651.000000
25%    98033.000000    47.471100    -122.328000    1490.000000
5100.000000
50%    98065.000000    47.571800    -122.231000    1840.000000
7620.000000
75%    98118.000000    47.678000    -122.125000    2360.000000
10083.000000
max     98199.000000    47.777600    -121.315000    6210.000000
871200.000000

```

```

.....Head:.....

```

	id	date	price	bedrooms	bathrooms	sqft_living
\						
0	7129300520	10/13/2014	221900.0	3	1.00	1180
1	6414100192	12/9/2014	538000.0	3	2.25	2570
2	5631500400	2/25/2015	180000.0	2	1.00	770
3	2487200875	12/9/2014	604000.0	4	3.00	1960
4	1954400510	2/18/2015	510000.0	3	2.00	1680

	sqft_lot	floors	waterfront	view	...	grade	sqft_above	\
0	5650	1.0	NaN	NONE	...	7 Average	1180	
1	7242	2.0	NO	NONE	...	7 Average	2170	
2	10000	1.0	NO	NONE	...	6 Low Average	770	
3	5000	1.0	NO	NONE	...	7 Average	1050	
4	8080	1.0	NO	NONE	...	8 Good	1680	

	sqft_basement	yr_built	yr_renovated	zipcode	lat	long	\
0	0.0	1955	0.0	98178	47.5112	-122.257	
1	400.0	1951	1991.0	98125	47.7210	-122.319	
2	0.0	1933	NaN	98028	47.7379	-122.233	
3	910.0	1965	0.0	98136	47.5208	-122.393	
4	0.0	1987	0.0	98074	47.6168	-122.045	

	sqft_living15	sqft_lot15
0	1340	5650
1	1690	7639
2	2720	8062
3	1360	5000
4	1800	7503

```

[5 rows x 21 columns]

```

## Null values

Looking at the information above we can see only three columns have missing values, that is; "waterfront", "view" and "yr\_renovated". Every house has its own unique features and not all are the same. Some houses contain certain features while others lack them. Since this is real world data, we can account for missing values in "waterfront" and "view" columns by saying not all houses are build the same and those lacking the two features have caused our data on the two columns to be inconsitent with the rest of the other columns. The "yr\_renovated" column can also be accounted for by saying not all houses undergo renovation. Houses build earlier might need renovation but recent houses do not require renovation hence the missing values in the column

```
# Using mode to impute missing values
# Python function to impute missing values

def replace_missing_with_mode(data, column_name):
    mode_value = data[column_name].mode().iloc[0]
    data[column_name].fillna(mode_value, inplace =True)

# columns to be imputed
replace_missing_with_mode(data, 'view')
replace_missing_with_mode(data, 'waterfront')

# Changing our date from object to datetime data type
data['date'] = pd.to_datetime(data['date'])
```

After checking for null values, we check for any duplicated values in the data.

```
# Checking for duplicated values in our data
data.duplicated().sum()

0

# drop the rows in sqft_basement with a '?'
data= data.drop(data[data.sqft_basement == '?'].index)
```

Creating a new column 'Grade\_1' that stores our new 'grade' column after getting rid of the string 'grade' and converting it to a numeric datatype

```
data["Grade_1"] = data["grade"].str.split().apply(lambda x: x[0])
# Convert the Grade1 column to an integer.
data["Grade_1"] = pd.to_numeric(data["Grade_1"])
```

## Converting the categorical columns to numerical data types

We are converting the following categorical data "Waterfront", "View" and "grade" into numerical data types.

```
data['view_1'] = data['view'].replace({'NONE': 0, 'FAIR':1, 'AVERAGE':
2, 'GOOD':3, 'EXCELLENT':4})
data['waterfront_1'] = data['waterfront'].replace({'YES': 0, 'NO':1})
data['condition1'] = data['condition'].replace({'Poor': 0,
'Fair':1, 'Average':2, 'Good':3, 'Very Good':4})
```

Since we have already replaced the strings in our categorical data with numeric values we can drop the original columns

```
data.drop(columns = ['waterfront', 'view', 'grade', 'condition', ],
inplace= True)
```

```
data.dtypes
```

```
id                int64
date              datetime64[ns]
price             float64
bedrooms          int64
bathrooms         float64
sqft_living       int64
sqft_lot          int64
floors            float64
sqft_above        int64
sqft_basement     object
yr_built          int64
yr_renovated      float64
zipcode           int64
lat              float64
long             float64
sqft_living15     int64
sqft_lot15        int64
Grade_1           int64
view_1           int64
waterfront_1      int64
condition1        int64
dtype: object
```

Delivery locations (zip codes) data

```
# Create a new dataframe with two columns for zip codes and cities
new_areas = pd.DataFrame(columns=['Zip Code', 'City'])

# Iterate over the original dataframe and extract zip codes and cities
for i in range(len(areas.columns)):
    # Skip the columns that are not zip codes
    if i % 2 != 0:
        continue
    # Extract the zip codes and cities from each pair of columns
```

```

zip_codes = areas.iloc[:, i]
cities = areas.iloc[:, i + 1]
# Append the zip codes and cities to the new dataframe
new_areas = new_areas.append(pd.DataFrame({'Zip Code': zip_codes,
'City': cities})), ignore_index=True)

```

```

# Print the new dataframe
print(new_areas)

```

	Zip Code	City
0	98001	Algona
1	98001	Auburn
2	98001	Federal Way
3	98002	Auburn
4	98003	Federal Way
...	...	...
447	NaN	NaN
448	NaN	NaN
449	NaN	NaN
450	NaN	NaN
451	NaN	NaN

[452 rows x 2 columns]

```

new_areas.isnull().sum()
new_areas.dropna()

```

	Zip Code	City
0	98001	Algona
1	98001	Auburn
2	98001	Federal Way
3	98002	Auburn
4	98003	Federal Way
...	...	...
404	98593	Vader
405	98595	Westport
406	98596	Chehalis
407	98596	Winlock
408	98597	Yelm

[409 rows x 2 columns]

```

#renaming to match our first data set
new_areas = new_areas.rename(columns={"Zip Code": "zipcode"})
new_areas

```

	zipcode	City
0	98001	Algona
1	98001	Auburn
2	98001	Federal Way
3	98002	Auburn

```

4      98003  Federal Way
..      ...      ...
447     NaN      NaN
448     NaN      NaN
449     NaN      NaN
450     NaN      NaN
451     NaN      NaN

```

```
[452 rows x 2 columns]
```

```
#merging our data sets
```

```

new_areas['zipcode'] = new_areas['zipcode'].astype(str)
data['zipcode'] = new_areas['zipcode'].astype(str)

```

```

merged_data = pd.merge(new_areas, data , on='zipcode')
merged_data

```

	zipcode	City	id	date	price	bedrooms
bathrooms \						
0	98001	Algona	7129300520	10/13/2014	221900.0	3
1.00						
1	98001	Algona	6414100192	12/9/2014	538000.0	3
2.25						
2	98001	Algona	5631500400	2/25/2015	180000.0	2
1.00						
3	98001	Auburn	7129300520	10/13/2014	221900.0	3
1.00						
4	98001	Auburn	6414100192	12/9/2014	538000.0	3
2.25						
...	...	...	...	...	...	...
...						
2405	nan	NaN	1049010390	3/19/2015	505000.0	3
2.00						
2406	nan	NaN	7905370390	10/9/2014	475000.0	5
2.50						
2407	nan	NaN	4140090240	11/5/2014	520000.0	3
2.25						
2408	nan	NaN	4055700030	5/2/2015	1450000.0	3
4.50						
2409	nan	NaN	3775300030	12/31/2014	333500.0	3
1.75						

	sqft_living	sqft_lot	floors	...	yr_built	yr_renovated
lat \						
0	1180	5650	1.0	...	1955	0.0
47.5112						
1	2570	7242	2.0	...	1951	1991.0
47.7210						
2	770	10000	1.0	...	1933	NaN



47.7379						
3	1180	5650	1.0	...	1955	0.0
47.5112						
4	2570	7242	2.0	...	1951	1991.0
47.7210						
...	...	...	...	...	...	...
..						
2405	1260	5460	1.0	...	1972	0.0
47.7355						
2406	2340	7200	1.0	...	1975	0.0
47.7206						
2407	2590	9263	1.0	...	1977	0.0
47.7691						
2408	3970	24920	2.0	...	1977	NaN
47.7183						
2409	1220	9732	1.0	...	1965	0.0
47.7736						

	long	sqft_living15	sqft_lot15	Grade_1	view_1
waterfront_1	\				
0	-122.257	1340	5650	7	0.0
NaN					
1	-122.319	1690	7639	7	0.0
1.0					
2	-122.233	2720	8062	6	0.0
1.0					
3	-122.257	1340	5650	7	0.0
NaN					
4	-122.319	1690	7639	7	0.0
1.0					
...	...	...	...	...	...
.					
2405	-122.180	1510	5460	7	0.0
1.0					
2406	-122.211	1930	7221	7	0.0
1.0					
2407	-122.262	2580	9450	8	0.0
1.0					
2408	-122.258	2610	13838	10	2.0
1.0					
2409	-122.214	1630	10007	7	0.0
1.0					

	condition1
0	2
1	2
2	2
3	2
4	2

```
...      ...
2405      2
2406      2
2407      4
2408      2
2409      2
```

```
[2410 rows x 22 columns]
```

```
merged_data.isnull().sum()
```

```
#dropping the rows with null values
```

```
new_data = merged_data.dropna()
```

```
def create_bar_graph(data):
```

```
    # Count the occurrences of each city
```

```
    city_counts = data['City'].value_counts()
```

```
    # Select the top ten cities
```

```
    top_cities = city_counts.head(10)
```

```
    # Create a bar graph
```

```
    plt.bar(top_cities.index, top_cities.values)
```

```
    plt.xlabel('City')
```

```
    plt.ylabel('Number of Houses')
```

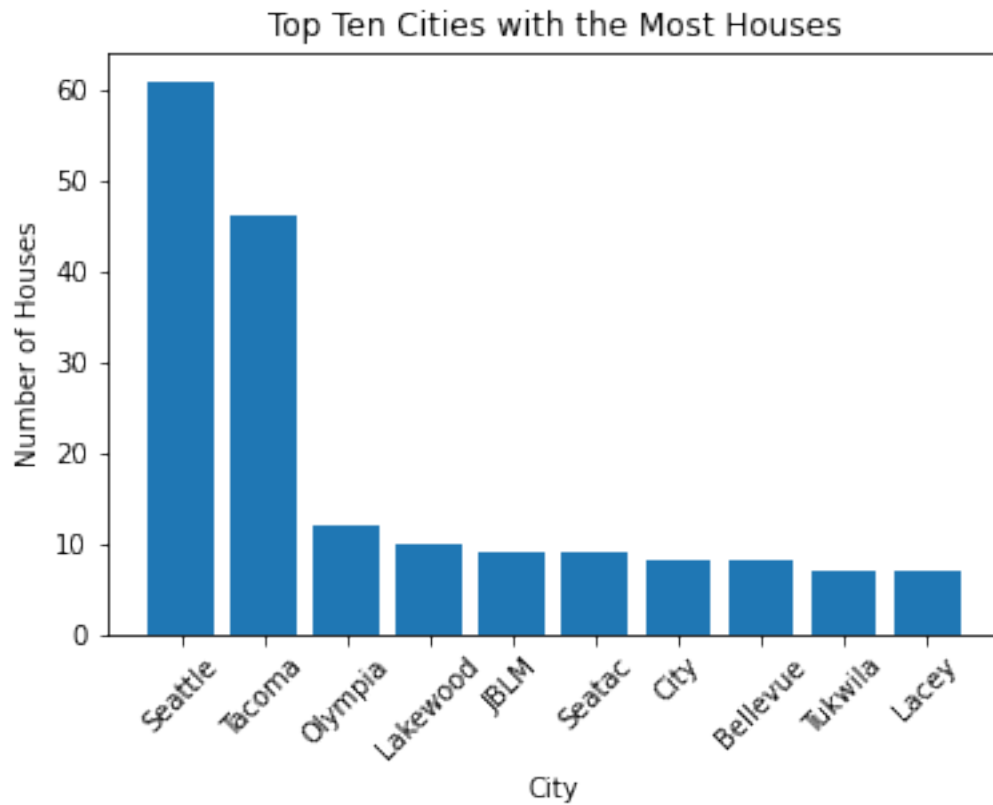
```
    plt.title('Top Ten Cities with the Most Houses')
```

```
    plt.xticks(rotation=45)
```


```
    plt.show()
```

```
# Example usage:
```

```
create_bar_graph(new_data)
```



```
city_prices = new_data.groupby('City')['price'].mean()
# Select the top ten cities with the highest mean prices
top_ten_cities = city_prices.nlargest(10)
# Create a bar graph
top_ten_cities.plot(kind='bar', xlabel='City', ylabel='Price',
title='Top Ten Cities with Highest Prices')
# Show the plot
plt.show()
```



## Data visualization

Now we have checked for abnormalities in the data, we can go ahead and plot the data to explore the distribution, relationships and patterns in the data. This will also help us in identifying outliers and trends.

```
y = data["price"]
X = data.drop("price", axis = 1)

def scatter_plots(y, X):
    plots = X.shape[1]
    cols = 4
    rows = (plots + cols - 1) // cols

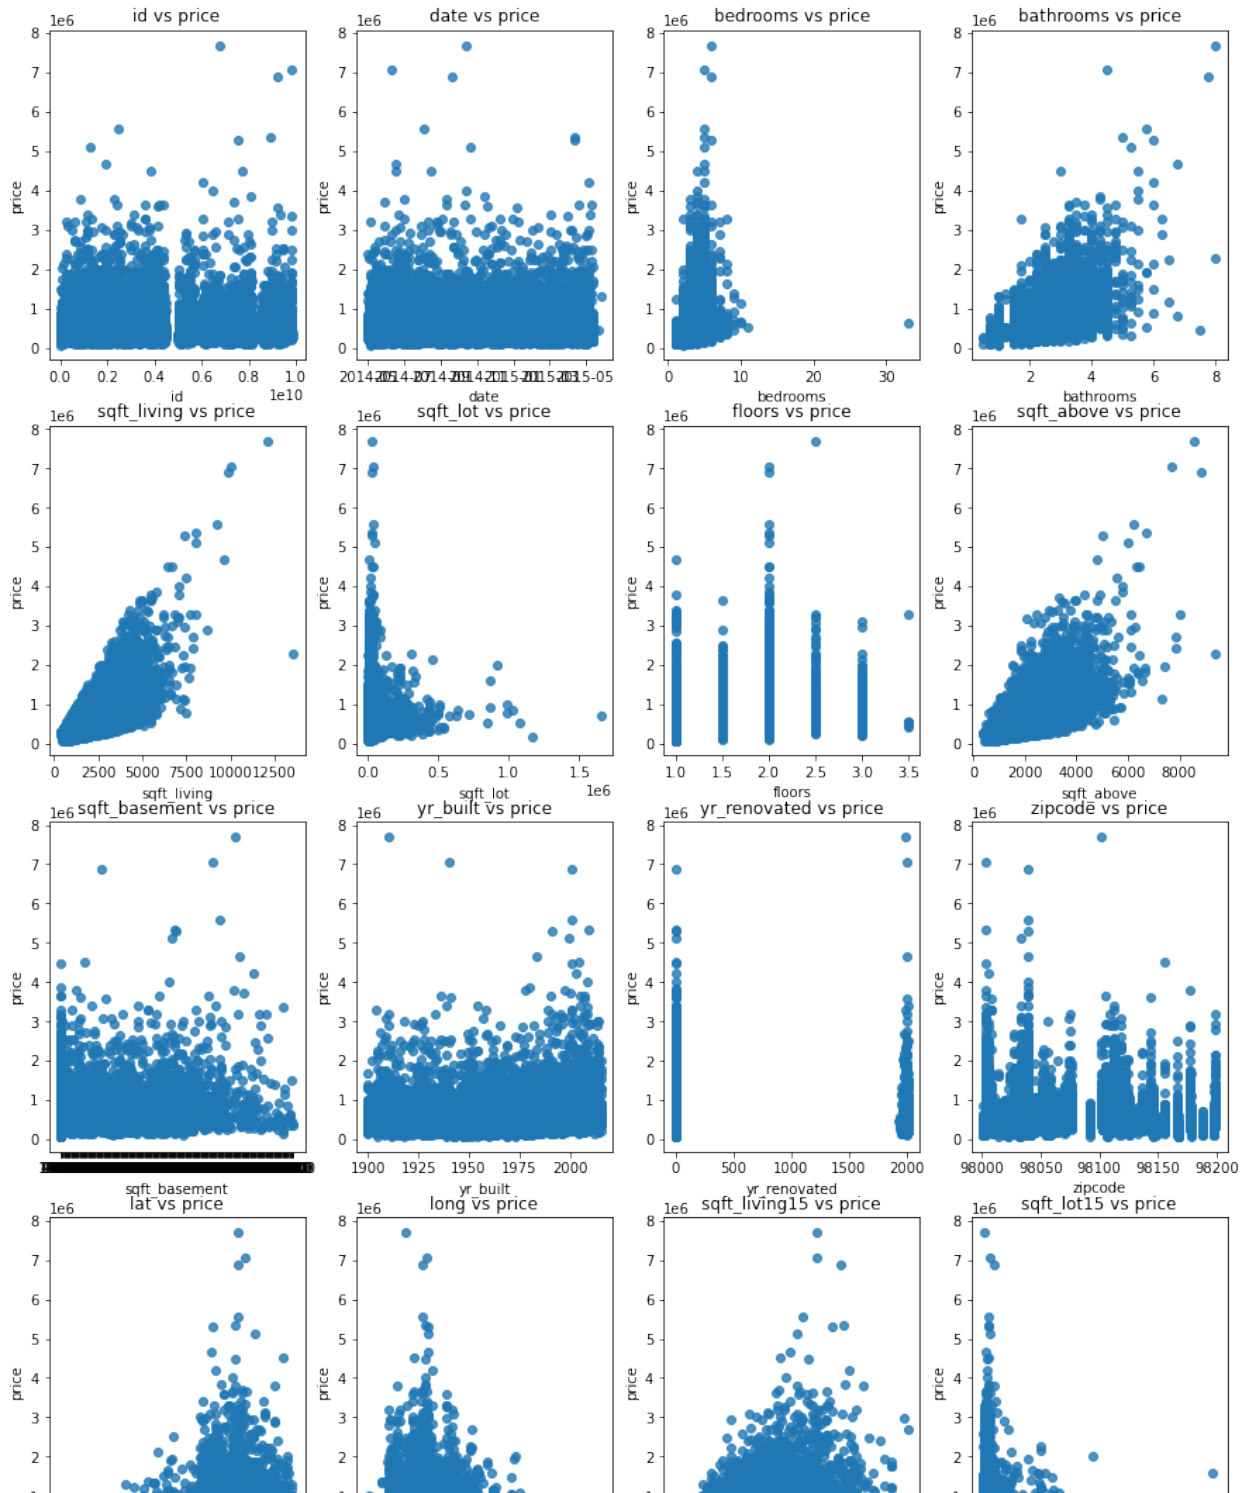
    fig, axes = plt.subplots(rows, cols, figsize=(15, 5 * rows))
    fig.suptitle(f"Scatter plot of Independent variables vs {y.name}")

    for i, ax in enumerate(axes.flat):
        if i < plots:
            x_col_name = X.columns[i]
            ax.scatter(X.iloc[:, i], y, alpha=0.8)
```

```
ax.set_xlabel(x_col_name)
ax.set_ylabel(y.name)
ax.set_title(f"{x_col_name} vs {y.name}")
```

```
# Run the function
scatter_plots(y, X)
```

Scatter plot of Independent variables vs price




## ## Relationship between our independent variables and the dependent variable("price")

```
#correlation
columns_to_test = data.columns

# computing the correlation matrix
correlation_matrix = data[columns_to_test].corr()
matrix = np.triu(np.ones_like(correlation_matrix, dtype = bool))
one_sided_correlation = correlation_matrix.mask(matrix)

# using heatmap to visualize the correlation
plt.figure(figsize=(15, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm',
linewidths=0.5, mask = matrix)
plt.title(f'Correlation Matrix')
plt.show()
```



## Correlation of our columns against the target("price")

```
def correlation(df):
    return data.corr()['price'].sort_values()

correlation(data)
```

```
waterfront_1    -0.265969
zipcode         -0.053166
id              -0.015796
long            0.022101
condition1      0.035290
yr_built        0.054459
sqft_lot15      0.083192
sqft_lot        0.087937
yr_renovated    0.128227
floors          0.256355
lat             0.306507
bedrooms        0.309204
view_1          0.394885
bathrooms       0.525889
sqft_living15   0.586415
sqft_above      0.605143
Grade_1         0.667738
sqft_living     0.702328
price           1.000000
Name: price, dtype: float64
```

## Analyzing price and location


```
def create_bar_graph(data):
    # Count the occurrences of each city
    city_counts = data['City'].value_counts()

    # Select the top ten cities
    top_cities = city_counts.head(10)


    # Create a bar graph
    plt.bar(top_cities.index, top_cities.values)
    plt.xlabel('City')
    plt.ylabel('Number of Houses')
    plt.title('Top Ten Cities with the Most Houses')
    plt.xticks(rotation=45)
    plt.show()

# Example usage:
create_bar_graph(new_data)
```





```
city_prices = new_data.groupby('City')['price'].mean()
# Select the top ten cities with the highest mean prices
top_ten_cities = city_prices.nlargest(10)
# Create a bar graph
top_ten_cities.plot(kind='bar', xlabel='City', ylabel='Price',
title='Top Ten Cities with Highest Prices')
# Show the plot
plt.show()
```



## Analyzing seasonal trends in prices

```
# Creating a function to map months to seasons
def get_season(date):
    if date.month in [3,4,5]:
        return 'Spring'
    elif date.month in [6,7,8]:
        return 'Summer'
    elif date.month in [9,10,11]:
        return 'Autumn'
    else:
        return 'Winter'

# Applying the function to the 'date' column to create a 'season'
column
data['season'] = data['date'].apply(get_season)
data[['date', 'season']]
```

	date	season
0	2014-10-13	Autumn
1	2014-12-09	Winter
2	2015-02-25	Winter
3	2014-12-09	Winter

```


4      2015-02-18  Winter
...
21592 2014-05-21  Spring
21593 2015-02-23  Winter
21594 2014-06-23  Summer
21595 2015-01-16  Winter
21596 2014-10-15  Autumn

[21597 rows x 2 columns]

```


Creating a boxplot of price segmented by season to view differences in price distribution by season.

```
sns.boxplot(x='season', y='price', data=data);
```



Making a timeseries plot of price over time, colored by season to see seasonal patterns.

```
data.set_index('date').groupby('season')['price'].plot();
```



Calculating summary statistics (mean, median, std dev) for price grouped by season to quantify differences.

```
data.groupby('season')['price'].agg([np.mean, np.median, np.std])
```

	mean	median	std
season			
Autumn	531276.474881	443725.0	378513.665722
Spring	552782.763271	465000.0	367075.050556
Summer	546719.464286	455000.0	368925.606702
Winter	519613.645467	430000.0	348171.543129

```
# Extract price by season into separate dataframes
```

```
spring = data[data['season'] == 'Spring']['price']
```

```
summer = data[data['season'] == 'Summer']['price']
```

```
fall = data[data['season'] == 'Autumn']['price']
```

```
winter = data[data['season'] == 'Winter']['price']
```

```
# Perform ANOVA test
```

```
f_val, p_val = stats.f_oneway(spring, summer, fall, winter)
```

```
print(f_val, p_val)
```

```
# Interpret results
```

```
alpha = 0.05
```

```
if p_val < alpha:
```

```
    print("We reject the null hypothesis")
```

```
    print("There is a statistically significant difference in price by
```

```
season")
else:
    print("We fail to reject the null hypothesis")
    print("There is no statistically significant difference in price by
season")
```

8.082642416374668 2.2312373653979034e-05

We reject the null hypothesis

There is a statistically significant difference in price by season

In this case; these are the hypotheses.

Null hypothesis:

There is no difference in the mean price across the seasons. The season has no effect on price.

$H_0: \mu_{\text{spring}} = \mu_{\text{summer}} = \mu_{\text{fall}} = \mu_{\text{winter}}$

Alternative hypothesis:

There is a difference in mean price for at least one season compared to the others. The season has an effect on price.

$H_1: \text{At least one } \mu_{\text{season}} \neq \mu_{\text{other seasons}}$

Where  $\mu_{\text{season}}$  is the population mean price for that season.

So in summary: Null hypothesis ( $H_0$ ): The seasons all have an equal effect on mean price (no difference). Alternative hypothesis ( $H_1$ ): At least two seasons have a statistically significant difference in their effect on mean price. If we reject  $H_0$  based on a small ANOVA p-value, we would conclude there is a significant difference in price by season. Failing to reject  $H_0$  means we cannot say there is a seasonal effect.

## Linear Regression

Since we now have a better understanding of the correlation between our target("price") and our features("independent variables"), we proceed to building regression models to further understand the magnitude our features have on price and predict whether this model can give us accurate house prices when fitted with the said features. We will explore a few features from our data set which we have investigated and come to a conclusion that they have a significance on our target. Steps involved here are as follows;

1. Feature Selection
2. Model Selection
3. Model training
4. Model evaluation
5. Model interpretation
6. Model validation and testing
7. \*\*feature engineering

## 1.Feature selection

Here we just choose the endogenous and exogenous variables. First we will select for the baseline model then after we shall add features to see if it improves and evaluate it.


```
y= data["price"]  
X = data["sqft_living"]
```

## 2. Building a baseline model

Creating a baseline model for our regression model

First we visualize the target(price) column in order to understand the distribution

```
sns.histplot(data["price"], color="maroon", kde=True)  
plt.show();
```



The distribution of our data seems to have a longer right tail than the left tail. This indicates a positive skewness in our target meaning the mean is greater than the median. This may have impact on our model since linear regression, assume that the target variable follows a normal distribution, so significant skewness can be problematic.

### Splitting our data into train and test sets

First we split the data into test set and training set. We will use the "train\_test\_split" function from scikit-learn library to split our data.

```

#importing scikit-learn library
from sklearn.model_selection import train_test_split
#defining a function for splitting data into train and test sets
def split(X, y):
    X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=0)
    return X_train, X_test, y_train, y_test

#splitting the data by calling the function
X_train, X_test, y_train, y_test = split(X,y)

```

In the train\_test\_split function we have is used above, we have split 80% of the data into training set and 20% of the data into test set.

In the next cell we are going to build a baseline model using the y\_train and x\_train variables. To do that we import statsmodels library which is a powerful library for statistical modelling and is similar to the sci-kit learn module

```

#importing statsmodels
import statsmodels.api as sm

#function to create models and print the summary
model = sm.OLS(y_train, sm.add_constant(X_train))
results = model.fit()
results.summary()

```

```

<class 'statsmodels.iolib.summary.Summary'>
"""
                                OLS Regression Results
=====
=====
Dep. Variable:                  price    R-squared:
0.495
Model:                          OLS      Adj. R-squared:
0.495
Method:                        Least Squares    F-statistic:
1.655e+04
Date:                          Thu, 26 Oct 2023    Prob (F-statistic):
0.00
Time:                          01:44:37    Log-Likelihood:
2.3497e+05
No. Observations:              16914    AIC:
4.700e+05
Df Residuals:                  16912    BIC:
4.700e+05
Df Model:                      1
Covariance Type:               nonrobust

```

```

=====
=====
              coef      std err          t      P>|t|      [0.025
0.975]
-----
-----
const      -4.448e+04    4960.688      -8.966      0.000    -5.42e+04
-3.48e+04
sqft_living    280.1781      2.178     128.639      0.000      275.909
284.447
=====
=====
Omnibus:                11863.674    Durbin-Watson:
2.008
Prob(Omnibus):          0.000    Jarque-Bera (JB):
453836.272
Skew:                   2.904    Prob(JB):
0.00
Kurtosis:               27.703    Cond. No.
5.62e+03
=====
=====

```

#### Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 5.62e+03. This might indicate that there are strong multicollinearity or other numerical problems.

"""

```
import scipy.stats as stats
```


```
# Generate a Q-Q plot
```

```
sm.qqplot(data["price"], line='s', fit=True, dist=stats.norm, loc=0,
scale=1)
```

```
plt.title('Price')
```

```
plt.show()
```





Wonderful! We now have our baseline model and from it we can interpret its metrics

## Interpretation

Looking at the summary above, we can see that the regression line we found was

$$\text{price} = \text{sqft\_living}285.8630 - 43,990$$

The model is statistically significant overall, with an F-statistic p-value well below 0.05. The model explains about 50% of the variance in price. The model `const` and `sqft_living` coefficients are both statistically significant, with t-statistic p-values well below 0.05. If a house had 0 sqft of living, we would expect price to be about -43,990 dollars. For each increase of 1 sqft in the living, we see an associated increase in price of about 280 dollars.

We now have our baseline model which we created using the train sets obtained from splitting our data. To analyze our model further we will predict our target("price") using the trained model then compare metrics to determine if our model is efficient or we need to adjust it. We also check if our model is under fitted or over fitted. To perform comparison we will import another module from sci-kit learn

```
#predicting the dependent variable
X_test = sm.add_constant(X_test)
y_pred = results.predict(X_test)

#defining a function to calculate metrics for our prediction model
def calculate_regression_metrics(y_test, y_pred):
    metrics = {}
```

```

# Calculate Mean Absolute Error (MAE)
metrics['MAE'] = mean_absolute_error(y_test, y_pred)

# Calculate Mean Squared Error (MSE)
metrics['MSE'] = mean_squared_error(y_test, y_pred)

# Calculate R-squared (coefficient of determination)
metrics['R-squared'] = r2_score(y_test, y_pred)

return metrics

metrics = calculate_regression_metrics(y_test, y_pred)
print(metrics)

{'MAE': 177033.87791345723, 'MSE': 70202650031.35713, 'R-squared':
0.48779757847019767}

```

## Comparison of our baseline model against the prediction model

### R-squared

Comparing the r-squared values from the test model and the prediction model we can note a difference. Our train model has a R-squared value of almost 50% while our test model has a value of almost 48%. This means our baseline model is good ...


### Residuals

Residuals are the difference between the true values and the values predicted by our model. We visualize to understand the distribution and also check if it meets the assumption of linearity; that is normal distribution

```


# Generate a Q-Q plot
sm.qqplot(results.resid, line='s', fit=True, dist=stats.norm, loc=0,
scale=1)
plt.title('Q-Q Plot of baseline residuals')
plt.show()

```



Quantile-quantile plots are used to assess whether a dataset follows a specific theoretical normal distribution. The visual above shows our model residuals almost follow a straight line but then curves at some point. The skewness or outliers in our target might be the cause of this.

```
#ploting histogram to show residuals distribution  
sns.histplot(results.resid, bins = "auto", kde = True)  
plt.title("Residual distribtution")  
plt.show();
```




From the above plot we can see that our baseline model residuals have a normal distribution. This meets one of the linearity assumptions of linear regression. Linear models make key assumptions one of which is that the errors or residuals follow a normal distribution. The normality assumption is essential for valid statistical inferences and hypothesis testing. Non-normal residuals can lead to biased parameter estimates, incorrect p-values, and unreliable confidence intervals.

## Target transformation

The non-normal distribution shown by the residuals (where it curves above) can be accounted for by the skewness of our target. Transforming the target variable can be an effective approach to make the data more closely approximate a normal distribution. By transforming the target variable, you aim to reduce skewness and make the data more symmetric, thus bringing it closer to a normal distribution. This, in turn, helps the residuals conform more closely to the normality assumption, which is crucial for the validity of the model. We shall log transform our target and see if our residuals will follow a normal distribution.

```
#log transforming the target
y_log = np.log(data["price"])
# visualizing to see its distribution
sns.histplot(y_log, bins = "auto", kde = True)
plt.title("price_log distribution")
plt.show();
```



Wonderful! Our target after transformation seems to follow an almost normal distribution. Next we create a model for the transformed target, then we shall visualize the residuals once again to see their distribution. We will also check if it improves our model or not.

## Log transformed target model

```
#splitting data into test and train test
split(X, y_log)

##creating a model
log_model = sm.OLS(y_train, sm.add_constant(X_train))
log_results = log_model.fit()
log_results.summary()
```

```
<class 'statsmodels.iolib.summary.Summary'>
"""
                                OLS Regression Results

=====
=====
Dep. Variable:                  price    R-squared:
0.484
Model:                            OLS    Adj. R-squared:
0.484
Method:                 Least Squares    F-statistic:
1.586e+04
Date:                   Thu, 26 Oct 2023    Prob (F-statistic):
0.00
```

```

Time:                                01:45:00   Log-Likelihood:
-7525.5
No. Observations:                    16914   AIC:
1.506e+04
Df Residuals:                        16912   BIC:
1.507e+04
Df Model:                            1

Covariance Type:                    nonrobust

=====
=====
              coef      std err          t      P>|t|      [0.025
0.975]
-----
-----
const          12.2207      0.007    1704.768      0.000      12.207
12.235
sqft_living     0.0004    3.15e-06    125.933      0.000      0.000
0.000
=====
=====
Omnibus:                3.650   Durbin-Watson:
1.980
Prob(Omnibus):          0.161   Jarque-Bera (JB):
3.629
Skew:                   0.035   Prob(JB):
0.163
Kurtosis:               3.015   Cond. No.
5.62e+03
=====
=====

Notes:
[1] Standard Errors assume that the covariance matrix of the errors is
correctly specified.
[2] The condition number is large, 5.62e+03. This might indicate that
there are
strong multicollinearity or other numerical problems.
"""

```

## Interpreting the results of the model

From the summary above we can see that the regression line we found was

price = sqft\_living0.0004 - 12.2207

The model is statistically significant overall, with an F-statistic p-value well below 0.05. The model explains about 48% of the variance in price. The model `const` and `sqft_living` coefficients are both statistically significant, with t-statistic p-values well below 0.05. If a house

had 0 sqft of living , we would expect price to be about -12.2207 dollars. For each increase of 1 sqft in the living , we see an associated increase in price of about 0.0004 dollars.

## Prediction of our model


```
X_test = sm.add_constant(X_test)
y_pred = results.predict(X_test)

calculate_regression_metrics(y_test, y_pred)

{'MAE': 536828.527684868,
 'MSE': 352139274884.2067,
 'R-squared': -1250053982156.927}
```

Visualizing the transformed target residuals

```
# Generate a Q-Q plot
sm.qqplot(log_results.resid, line='s', fit=True, dist=stats.norm,
loc=0, scale=1)
plt.title('Q-Q Plot of our prediction model residuals')
plt.show()
```



The residuals of the log transformed target now follow a straight line compared to our first model. This is a good indication since it now meets the assumption of normal distribution.

## Evaluation between the two baseline models

The first baseline model is much better than the second baseline model since it explains almost 50% in variance of the target variable compared to the second one which explains almost 48%. The second model prediction is also very poor since it has a negative R-squared value which indicates that the regression model's fit to the data is worse the model is not explaining any of the variance in the dependent variable, and it might be a poor fit for the data.

## 2nd Model

In this model we are going to improve our first baseline model which we didn't transform by adding more features and see if our model is going to improve.

Selecting y and X variables for our model

```
#defining variables to be used in our second model
drop = data.drop(['id', 'price',
                  'date', 'sqft_lot', 'floors', 'sqft_basement', 'yr_renovated',
                  'zipcode',
                  'lat', 'long', 'sqft_living15', 'sqft_lot15', "sqft_above"], axis= 1)
X_sec = drop
y_sec= data["price"]
X_sec
```

	bedrooms	bathrooms	sqft_living	sqft_above	yr_built	Grade_1
\						
0	3	1.00	1180	1180	1955	7
1	3	2.25	2570	2170	1951	7
2	2	1.00	770	770	1933	6
3	4	3.00	1960	1050	1965	7
4	3	2.00	1680	1680	1987	8
...	...	...	...	...	...	...
21592	3	2.50	1530	1530	2009	8
21593	4	2.50	2310	2310	2014	8
21594	2	0.75	1020	1020	2009	7
21595	3	2.50	1600	1600	2004	8
21596	2	0.75	1020	1020	2008	7
	view_1	waterfront_1	condition1			
0	0	1	2			



1	0	1	2
2	0	1	2
3	0	1	4
4	0	1	2
...	...	...	...
21592	0	1	2
21593	0	1	2
21594	0	1	2
21595	0	1	2
21596	0	1	2

[21143 rows x 9 columns]

In this model we have standardized the X variables (features) to have a mean of 0 and a standard deviation of 1 is known as standardization or z-score normalization. Scaling the features to have a standard deviation of 1 ensures that the features have the same variance, which can be important for modelling algorithms. Standardizing the features can help make the data closer to a normal distribution, which can improve the model's performance for such models.

```
ss = StandardScaler()
X1_scaled = ss.fit_transform(X_sec)
```

splitting data into train and test

```
split(X1_scaled, y_sec)
```

Building multiple linear model

```
model = sm.OLS(y_train, sm.add_constant(X_train))
results = model.fit()
results.summary()
```

```
<class 'statsmodels.iolib.summary.Summary'>
"""
```

#### OLS Regression Results

```
=====
=====
Dep. Variable:                price    R-squared:
0.650
Model:                        OLS      Adj. R-squared:
0.649
Method:                      Least Squares    F-statistic:
3481.
Date:                        Thu, 26 Oct 2023    Prob (F-statistic):
0.00
Time:                        01:45:16    Log-Likelihood:
2.3188e+05
```

No. Observations: 16914 AIC:  
 4.638e+05  
 Df Residuals: 16904 BIC:  
 4.639e+05  
 Df Model: 9

Covariance Type: nonrobust

	coef	std err	t	P> t	[0.025	0.975]
const	5.393e+05	1673.474	322.268	0.000	5.36e+05	5.43e+05
x1	-3.469e+04	2111.891	-16.428	0.000	-3.88e+04	-3.06e+04
x2	4.268e+04	2889.968	14.767	0.000	3.7e+04	4.83e+04
x3	1.476e+05	4505.256	32.765	0.000	1.39e+05	1.56e+05
x4	4653.5416	3826.021	1.216	0.224	-2845.858	1.22e+04
x5	-1.049e+05	2221.989	-47.216	0.000	-1.09e+05	-1.01e+05
x6	1.491e+05	2882.270	51.727	0.000	1.43e+05	1.55e+05
x7	3.188e+04	1950.901	16.339	0.000	2.81e+04	3.57e+04
x8	-5.399e+04	1872.819	-28.826	0.000	-5.77e+04	-5.03e+04
x9	1.13e+04	1812.824	6.231	0.000	7742.283	1.48e+04
Omnibus:	12604.697	Durbin-Watson:				
2.006						
Prob(Omnibus):	0.000	Jarque-Bera (JB):				
826533.213						
Skew:	2.997	Prob(JB):				
0.00						
Kurtosis:	36.718	Cond. No.				
6.64						

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is

```
correctly specified.  
"""
```

Our model is good overall with an R-squared value of 65% meaning it explains 65 % of variance in our target variable. The p-values of our independent variables are below 0.05 meaning our model is statistically significant overall.

```
X_test = sm.add_constant(X_test)  
y_pred = results.predict(X_test)  
  
calculate_regression_metrics(y_test, y_pred)  
  
{ 'MAE': 141592.73032447265,  
  'MSE': 46826238256.493065,  
  'R-squared': 0.6583531730583653}
```

## Comparing the train and test model of our multiple linear model

The R-squared of our train model is 65% and that of our test model is also 65%. This can be an indication that our model is performing consistently between the training and test datasets. This can be a positive sign, suggesting that our model has not overfit the training data. This is a good model because the test model has not predicted more than what the train model predicted

metrics

## 3. Polynomial Transformation of features

We will use polynomial transformation to see if our model will improve or not.

```
# Polynomial transforming  
y_pol= data["price"]  
X_pol = X_sec  
X_pol
```

	bedrooms	bathrooms	sqft_living	sqft_above	yr_built	Grade_1
0	3	1.00	1180	1180	1955	7
1	3	2.25	2570	2170	1951	7
2	2	1.00	770	770	1933	6
3	4	3.00	1960	1050	1965	7
4	3	2.00	1680	1680	1987	8
...	...	...	...	...	...	...
21592	3	2.50	1530	1530	2009	8

21593	4	2.50	2310	2310	2014	8
21594	2	0.75	1020	1020	2009	7
21595	3	2.50	1600	1600	2004	8
21596	2	0.75	1020	1020	2008	7

	view_1	waterfront_1	condition1
0	0	1	2
1	0	1	2
2	0	1	2
3	0	1	4
4	0	1	2
...	...	...	...
21592	0	1	2
21593	0	1	2
21594	0	1	2
21595	0	1	2
21596	0	1	2

[21143 rows x 9 columns]

He we also standardized our data to have a mean of 0 and standard deviation of 1

```
ss = StandardScaler()
X1_scaled = ss.fit_transform(X_pol)
```

splitting the data

```
split(X1_scaled, y_sec)
```

Building the polynomial model

```
def build_polynomial_linear_model(X, y, degree=2):
    # Create polynomial features
    poly = PolynomialFeatures(degree=degree, include_bias=False)
    X_poly = poly.fit_transform(X)

    # Add a constant term (intercept)
    X_poly = (X_poly)

    # Build and fit a linear regression model using statsmodels
    model = sm.OLS(y, X_poly).fit()

    return model
```

```
# Build the polynomial linear model
model = build_polynomial_linear_model(X_train, y_train, degree=2)

# Print the summary of the model
model.summary()
```

```
<class 'statsmodels.iolib.summary.Summary'>
"""
```

### OLS Regression Results

```
=====
=====
```

Dep. Variable:	price	R-squared:	0.719
Model:	OLS	Adj. R-squared:	0.719
Method:	Least Squares	F-statistic:	815.8
Date:	Thu, 26 Oct 2023	Prob (F-statistic):	0.00
Time:	01:46:05	Log-Likelihood:	-2.3000e+05
No. Observations:	16914	AIC:	4.601e+05
Df Residuals:	16860	BIC:	4.605e+05
Df Model:	53		

Covariance Type: nonrobust

```
=====
=====
```

	coef	std err	t	P> t	[0.025	0.975]
x1	-1.334e+04	2241.953	-5.952	0.000	-1.77e+04	-8949.858
x2	2.851e+04	2972.717	9.591	0.000	2.27e+04	3.43e+04
x3	1.458e+05	5967.287	24.434	0.000	1.34e+05	1.57e+05
x4	-3.598e+04	4920.108	-7.312	0.000	-4.56e+04	-2.63e+04
x5	-7.013e+04	2713.742	-25.844	0.000	-7.55e+04	-6.48e+04

6.48e+04					
x6	1.373e+05	2775.927	49.475	0.000	1.32e+05
1.43e+05					
x7	1.559e+04	5007.623	3.114	0.002	5778.958
2.54e+04					
x8	5.948e+06	4.21e+04	141.228	0.000	5.87e+06
6.03e+06					
x9	1.931e+04	2446.360	7.894	0.000	1.45e+04
2.41e+04					
x10	638.3932	252.167	2.532	0.011	144.119
1132.668					
x11	3431.9165	2932.150	1.170	0.242	-2315.405
9179.238					
x12	-8264.1989	4494.629	-1.839	0.066	-1.71e+04
545.744					
x13	6834.6904	4042.086	1.691	0.091	-1088.221
1.48e+04					
x14	-3094.1989	2344.446	-1.320	0.187	-7689.558
1501.161					
x15	-3005.7492	3100.612	-0.969	0.332	-9083.274
3071.775					
x16	702.2192	2066.712	0.340	0.734	-3348.753
4753.192					
x17	-2103.9682	1962.403	-1.072	0.284	-5950.483
1742.547					
x18	-9.0281	2066.704	-0.004	0.997	-4059.984
4041.928					
x19	-6142.8575	2645.867	-2.322	0.020	-1.13e+04
-956.682					
x20	7820.0516	5934.096	1.318	0.188	-3811.398
1.95e+04					
x21	3564.2530	5237.666	0.681	0.496	-6702.122
1.38e+04					
x22	1.037e+04	3258.857	3.181	0.001	3979.567
1.68e+04					
x23	1.484e+04	4213.910	3.521	0.000	6579.119
2.31e+04					
x24	-591.2297	2533.067	-0.233	0.815	-5556.306
4373.847					
x25	-9710.6946	2166.193	-4.483	0.000	-1.4e+04 -
5464.729					
x26	-90.7363	2778.952	-0.033	0.974	-5537.773
5356.300					
x27	-6.847e+04	6956.386	-9.843	0.000	-8.21e+04 -
5.48e+04					
x28	1.324e+05	1.29e+04	10.263	0.000	1.07e+05
1.58e+05					
x29	-3.224e+04	5390.309	-5.980	0.000	-4.28e+04 -
2.17e+04					

x30 7.78e+04	6.559e+04	6218.224	10.548	0.000	5.34e+04	
x31 3180.173	-9808.4678	3381.602	-2.901	0.004	-1.64e+04	-
x32 1.18e+04	-1.649e+04	2407.166	-6.848	0.000	-2.12e+04	-
x33 2.03e+04	1.2e+04	4252.296	2.821	0.005	3660.688	
x34 4.68e+04	-6.102e+04	7239.149	-8.430	0.000	-7.52e+04	-
x35 1.09e+04	988.7095	5056.387	0.196	0.845	-8922.338	
x36 1.97e+04	-3.029e+04	5403.455	-5.606	0.000	-4.09e+04	-
x37 5051.385	-1.112e+04	3094.758	-3.592	0.000	-1.72e+04	-
x38 4589.683	-9445.1735	2477.159	-3.813	0.000	-1.43e+04	-
x39 6408.289	-1072.1197	3816.331	-0.281	0.779	-8552.529	
x40 2.71e+04	2.318e+04	1993.588	11.628	0.000	1.93e+04	
x41 -1.4e+04	-2.043e+04	3294.082	-6.202	0.000	-2.69e+04	
x42 7432.239	3191.8380	2163.355	1.475	0.140	-1048.563	
x43 2297.417	-7074.8195	2437.320	-2.903	0.004	-1.19e+04	-
x44 -402.971	-4555.7862	2118.670	-2.150	0.032	-8708.602	
x45 1.08e+04	5550.3683	2700.825	2.055	0.040	256.469	
x46 2.13e+04	1.62e+04	2608.060	6.212	0.000	1.11e+04	
x47 1.81e+04	1.406e+04	2078.162	6.766	0.000	9986.589	
x48 1.35e+04	7713.3382	2966.867	2.600	0.009	1897.967	
x49 7023.857	4060.2728	1511.952	2.685	0.007	1096.688	
x50 8643.404	4985.3976	1866.230	2.671	0.008	1327.391	
x51 7499.476	4054.5667	1757.513	2.307	0.021	609.657	
x52 5.03e+05	4.965e+05	3524.436	140.886	0.000	4.9e+05	
x53 3971.764	859.3544	1587.879	0.541	0.588	-2253.055	
x54	-889.9978	1311.948	-0.678	0.498	-3461.554	

```

1681.559
=====
=====
Omnibus:              7713.330    Durbin-Watson:
1.997
Prob(Omnibus):        0.000    Jarque-Bera (JB):
171655.002
Skew:                 1.678    Prob(JB):
0.00
Kurtosis:             18.242    Cond. No.
369.
=====
=====

Notes:
[1] Standard Errors assume that the covariance matrix of the errors is
correctly specified.
"""

```

## Interpreting the polynomial results

Our model is good overall with an R-squared value of 71% meaning it explains 71 % of variance in our target variable. The p-values of our independent variables are below 0.05 meaning our model is statistically significant overall.

coefficients and p\_values

```

model.params[:5]

x1      -1.334432e+04
x2       2.851020e+04
x3       1.458025e+05
x4      -3.597659e+04
x5      -7.013461e+04
x6       1.373397e+05
x7       1.559442e+04
x8       5.947704e+06
x9       1.931195e+04
x10      6.383932e+02
x11      3.431917e+03
x12     -8.264199e+03
x13      6.834690e+03
x14     -3.094199e+03
x15     -3.005749e+03
x16      7.022192e+02
x17     -2.103968e+03
x18     -9.028054e+00
x19     -6.142858e+03
x20      7.820052e+03

```



```
x21      3.564253e+03
x22      1.036727e+04
x23      1.483882e+04
x24     -5.912297e+02
x25     -9.710695e+03
x26     -9.073632e+01
x27     -6.847166e+04
x28      1.324299e+05
x29     -3.223643e+04
x30      6.558897e+04
x31     -9.808468e+03
x32     -1.648529e+04
x33      1.199563e+04
x34     -6.102379e+04
x35      9.887095e+02
x36     -3.029283e+04
x37     -1.111744e+04
x38     -9.445173e+03
x39     -1.072120e+03
x40      2.318242e+04
x41     -2.042966e+04
x42      3.191838e+03
x43     -7.074819e+03
x44     -4.555786e+03
x45      5.550368e+03
x46      1.620088e+04
x47      1.406000e+04
x48      7.713338e+03
x49      4.060273e+03
x50      4.985398e+03
x51      4.054567e+03
x52      4.965438e+05
x53      8.593544e+02
x54     -8.899978e+02
dtype: float64
```

```
model.pvalues[:5]
```

```
x1      2.699746e-09
x2      9.948197e-22
x3      1.335751e-129
x4      2.746822e-13
x5      1.813546e-144
x6      0.000000e+00
x7      1.847929e-03
x8      0.000000e+00
x9      3.101291e-15
x10     1.136250e-02
x11     2.418390e-01
x12     6.597936e-02
```

```
x13      9.087783e-02
x14      1.869198e-01
x15      3.323570e-01
x16      7.340295e-01
x17      2.836730e-01
x18      9.965146e-01
x19      2.026199e-02
x20      1.875829e-01
x21      4.961947e-01
x22      1.469021e-03
x23      4.304233e-04
x24      8.154500e-01
x25      7.414201e-06
x26      9.739531e-01
x27      8.466214e-23
x28      1.224674e-24
x29      2.270258e-09
x30      6.258948e-26
x31      3.730033e-03
x32      7.724659e-12
x33      4.793341e-03
x34      3.742282e-17
x35      8.449751e-01
x36      2.100622e-08
x37      3.286498e-04
x38      1.378315e-04
x39      7.787680e-01
x40      3.882083e-31
x41      5.707568e-10
x42      1.401206e-01
x43      3.704348e-03
x44      3.154532e-02
x45      3.988816e-02
x46      5.359265e-10
x47      1.371273e-11
x48      9.335244e-03
x49      7.250297e-03
x50      7.561412e-03
x51      2.106737e-02
x52      0.000000e+00
x53      5.883794e-01
x54      4.975410e-01
dtype: float64
```

Prediction metrics

```
X_test = sm.add_constant(X_test)
y_pred = results.predict(X_test)
```

```
calculate_regression_metrics(y_test, y_pred)
{'MAE': 141592.73032447265,
 'MSE': 46826238256.493065,
 'R-squared': 0.6583531730583653}
```

## Comparing the polynomial train and test model

On looking on the metrics of the two models we can say that our training data did well in training the model. We can see that the training model has a R-squared value of 65 while our training model has a R-squared of 71%. This shows our test model predicts well since it did not exceed the training model value and it also means that we did not overfit our data.